

```
from sklearn.datasets import load_breast_cancer
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

Loading a dataset about breast cancer, where we have information about tumors (like size, shape, and texture). The dataset also tells us whether each tumor is **benign** (not harmful) or **malignant** (harmful). To put this data into a table

```
# Load the dataset
breast_cancer_data = load_breast_cancer()

# Convert to a pandas DataFrame for better readability
data_df = pd.DataFrame(breast_cancer_data.data, columns=breast_cancer_data.feature_names)
target_df = pd.DataFrame(breast_cancer_data.target, columns=["target"])

# Concatenate feature and target DataFrames
full_data = pd.concat([data_df, target_df], axis=1)

# Display the first few rows of the DataFrame
print("Dataset Preview:")
print(full_data.head())
```

+ Code

+ Text



Dataset Preview:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	\
0	17.99	10.38	122.80	1001.0	0.11840	
1	20.57	17.77	132.90	1326.0	0.08474	
2	19.69	21.25	130.00	1203.0	0.10960	
3	11.42	20.38	77.58	386.1	0.14250	
4	20.29	14.34	135.10	1297.0	0.10030	

  

	mean compactness	mean concavity	mean concave points	mean symmetry	\
0	0.27760	0.3001	0.14710	0.2419	
1	0.07864	0.0869	0.07017	0.1812	
2	0.15990	0.1974	0.12790	0.2069	
3	0.28390	0.2414	0.10520	0.2597	
4	0.13280	0.1980	0.10430	0.1809	

  

	mean fractal dimension	...	worst texture	worst perimeter	worst area	\
0	0.07871	...	17.33	184.60	2019.0	
1	0.05667	...	23.41	158.80	1956.0	
2	0.05999	...	25.53	152.50	1709.0	
3	0.09744	...	26.50	98.87	567.7	
4	0.05883	...	16.67	152.20	1575.0	

  

	worst smoothness	worst compactness	worst concavity	worst concave points	\
0	0.1622	0.6656	0.7119	0.2654	
1	0.1238	0.1866	0.2416	0.1860	
2	0.1444	0.4245	0.4504	0.2430	
3	0.2098	0.8663	0.6869	0.2575	
4	0.1374	0.2050	0.4000	0.1625	

  

	worst symmetry	worst fractal dimension	target
0	0.4601	0.11890	0
1	0.2750	0.08902	0
2	0.3613	0.08758	0
3	0.6638	0.17300	0
4	0.2364	0.07678	0

[5 rows x 31 columns]

Double-click (or enter) to edit

```
# Summary of the dataset
print("\nSummary Statistics:")
print(data_df.describe())
```



Summary Statistics:

	mean radius	mean texture	mean perimeter	mean area	\
count	569.000000	569.000000	569.000000	569.000000	
mean	14.127292	19.289649	91.969033	654.889104	
std	3.524049	4.301036	24.298981	351.914129	
min	6.981000	9.710000	43.790000	143.500000	
25%	11.700000	16.170000	75.170000	420.300000	

50%	13.370000	18.840000	86.240000	551.100000
75%	15.780000	21.800000	104.100000	782.700000
max	28.110000	39.280000	188.500000	2501.000000

	mean smoothness	mean compactness	mean concavity	mean concave points \
count	569.000000	569.000000	569.000000	569.000000
mean	0.096360	0.104341	0.088799	0.048919
std	0.014064	0.052813	0.079720	0.038803
min	0.052630	0.019380	0.000000	0.000000
25%	0.086370	0.064920	0.029560	0.020310
50%	0.095870	0.092630	0.061540	0.033500
75%	0.105300	0.130400	0.130700	0.074000
max	0.163400	0.345400	0.426800	0.201200

	mean symmetry	mean fractal dimension	...	worst radius \
count	569.000000	569.000000	...	569.000000
mean	0.181162	0.062798	...	16.269190
std	0.027414	0.007060	...	4.833242
min	0.106000	0.049960	...	7.930000
25%	0.161900	0.057700	...	13.010000
50%	0.179200	0.061540	...	14.970000
75%	0.195700	0.066120	...	18.790000
max	0.304000	0.097440	...	36.040000

	worst texture	worst perimeter	worst area	worst smoothness \
count	569.000000	569.000000	569.000000	569.000000
mean	25.677223	107.261213	880.583128	0.132369
std	6.146258	33.602542	569.356993	0.022832
min	12.020000	50.410000	185.200000	0.071170
25%	21.080000	84.110000	515.300000	0.116600
50%	25.410000	97.660000	686.500000	0.131300
75%	29.720000	125.400000	1084.000000	0.146000
max	49.540000	251.200000	4254.000000	0.222600

	worst compactness	worst concavity	worst concave points \
count	569.000000	569.000000	569.000000
mean	0.254265	0.272188	0.114606
std	0.157336	0.208624	0.065732
min	0.027290	0.000000	0.000000
25%	0.147200	0.114500	0.064930
50%	0.211900	0.226700	0.099930
75%	0.339100	0.382900	0.161400
max	1.058000	1.252000	0.291000

	worst symmetry	worst fractal dimension
count	569.000000	569.000000
mean	0.290076	0.083946
std	0.061867	0.018061
min	0.156500	0.055040
25%	0.250400	0.071460

```
# Target class distribution
print("\nTarget Class Distribution:")
print(target_df['target'].value_counts())
```



```
Target Class Distribution:
target
1      357
0      212
Name: count, dtype: int64
```

## Handling Missing Values

```
# Check for missing values
print(full_data.isnull().sum())
```



```
mean radius      0
mean texture     0
mean perimeter   0
mean area        0
mean smoothness  0
mean compactness 0
mean concavity   0
mean concave points 0
mean symmetry    0
mean fractal dimension 0
radius error     0
texture error    0
perimeter error  0
```

```

area error          0
smoothness error    0
compactness error   0
concavity error     0
concave points error 0
symmetry error      0
fractal dimension error 0
worst radius        0
worst texture       0
worst perimeter     0
worst area          0
worst smoothness    0
worst compactness   0
worst concavity     0
worst concave points 0
worst symmetry      0
worst fractal dimension 0
target             0
dtype: int64

```

**Feature Scaling:** Performed using StandardScaler to ensure equal importance for all features in the model training process.

The features have different units and scales. Without scaling, the model could disproportionately emphasize larger-scale features.

StandardScaler standardizes the data, transforming it so that it has a mean of 0 and a standard deviation of 1.

Start coding or [generate](#) with AI.

```

# Feature Scaling: Standardize the features
scaler = StandardScaler()
data_scaled = scaler.fit_transform(data_df) # This scales the features

# Convert the scaled data back to a DataFrame for readability
data_scaled_df = pd.DataFrame(data_scaled, columns=breast_cancer_data.feature_names)

# Show the first few rows of the scaled data
print("\nScaled Data Preview:")
print(data_scaled_df.head())

```



Scaled Data Preview:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	\
0	1.097064	-2.073335	1.269934	0.984375	1.568466	
1	1.829821	-0.353632	1.685955	1.908708	-0.826962	
2	1.579888	0.456187	1.566503	1.558884	0.942210	
3	-0.768909	0.253732	-0.592687	-0.764464	3.283553	
4	1.750297	-1.151816	1.776573	1.826229	0.280372	

  

	mean compactness	mean concavity	mean concave points	mean symmetry	\
0	3.283515	2.652874	2.532475	2.217515	
1	-0.487072	-0.023846	0.548144	0.001392	
2	1.052926	1.363478	2.037231	0.939685	
3	3.402909	1.915897	1.451707	2.867383	
4	0.539340	1.371011	1.428493	-0.009560	

  

	mean fractal dimension	...	worst radius	worst texture	worst perimeter	\
0	2.255747	...	1.886690	-1.359293	2.303601	
1	-0.868652	...	1.805927	-0.369203	1.535126	
2	-0.398008	...	1.511870	-0.023974	1.347475	
3	4.910919	...	-0.281464	0.133984	-0.249939	
4	-0.562450	...	1.298575	-1.466770	1.338539	

  

	worst area	worst smoothness	worst compactness	worst concavity	\
0	2.001237	1.307686	2.616665	2.109526	
1	1.890489	-0.375612	-0.430444	-0.146749	
2	1.456285	0.527407	1.082932	0.854974	
3	-0.550021	3.394275	3.893397	1.989588	
4	1.220724	0.220556	-0.313395	0.613179	

  

	worst concave points	worst symmetry	worst fractal dimension
0	2.296076	2.750622	1.937015
1	1.087084	-0.243890	0.281190
2	1.955000	1.152255	0.201391
3	2.175786	6.046041	4.935010
4	0.729259	-0.868353	-0.397100

[5 rows x 30 columns]

**Dataset Split:** The dataset was split into training and testing sets to allow for model evaluation on unseen data.

```
# Split the dataset into train and test sets (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(data_df, target_df, test_size=0.2, random_state=42)
```

**Logistic Regression:** A simple linear classifier. Suitable for datasets with linear relationships.

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

```
# Initialize and train the model
log_reg = LogisticRegression(max_iter=10000)
log_reg.fit(X_train, y_train)
```

```
# Predict on the test data
y_pred_log_reg = log_reg.predict(X_test)
```

```
# Evaluate the model
log_reg_accuracy = accuracy_score(y_test, y_pred_log_reg)
print(f"Logistic Regression Accuracy: {log_reg_accuracy:.4f}")
```

```
→ /usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1408: DataConversionWarning: A column-vector y was passed when a 1d
y = column_or_1d(y, warn=True)
Logistic Regression Accuracy: 0.9561
```

**Decision Tree:** A non-linear classifier. Useful for capturing complex decision boundaries and interpretable models.

```
from sklearn.tree import DecisionTreeClassifier
```

```
# Initialize and train the model
dt_clf = DecisionTreeClassifier(random_state=42)
dt_clf.fit(X_train, y_train)
```

```
# Predict on the test data
y_pred_dt = dt_clf.predict(X_test)
```

```
# Evaluate the model
dt_accuracy = accuracy_score(y_test, y_pred_dt)
print(f"Decision Tree Accuracy: {dt_accuracy:.4f}")
```

```
→ Decision Tree Accuracy: 0.9474
```

**Random Forest:** An ensemble method combining multiple decision trees. Helps reduce overfitting and is robust to noise.

```
from sklearn.ensemble import RandomForestClassifier
```

```
# Initialize and train the model
rf_clf = RandomForestClassifier(random_state=42)
rf_clf.fit(X_train, y_train)
```

```
# Predict on the test data
y_pred_rf = rf_clf.predict(X_test)
```

```
# Evaluate the model
rf_accuracy = accuracy_score(y_test, y_pred_rf)
print(f"Random Forest Accuracy: {rf_accuracy:.4f}")
```

```
→ /usr/local/lib/python3.10/dist-packages/sklearn/base.py:1389: DataConversionWarning: A column-vector y was passed when a 1d array was ex
return fit_method(estimator, *args, **kwargs)
Random Forest Accuracy: 0.9649
```

**SVM:** Finds the optimal hyperplane that maximizes the margin between classes. Works well in high-dimensional spaces.

```
from sklearn.svm import SVC
```

```
# Initialize and train the model
svm_clf = SVC(kernel='linear', random_state=42)
```

```
svm_clf.fit(X_train, y_train)
```

```
# Predict on the test data
y_pred_svm = svm_clf.predict(X_test)
```

```
# Evaluate the model
svm_accuracy = accuracy_score(y_test, y_pred_svm)
print(f"SVM Accuracy: {svm_accuracy:.4f}")
```

```
↳ /usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1408: DataConversionWarning: A column-vector y was passed when a 1d
y = column_or_1d(y, warn=True)
SVM Accuracy: 0.9561
```

**k-NN:** A non-parametric method that classifies based on the majority of k nearest neighbors. Simple and effective for non-linear decision boundaries.

```
from sklearn.neighbors import KNeighborsClassifier
```

```
# Initialize and train the model
knn_clf = KNeighborsClassifier(n_neighbors=5)
knn_clf.fit(X_train, y_train)
```

```
# Predict on the test data
y_pred_knn = knn_clf.predict(X_test)
```

```
# Evaluate the model
knn_accuracy = accuracy_score(y_test, y_pred_knn)
print(f"k-NN Accuracy: {knn_accuracy:.4f}")
```

```
↳ k-NN Accuracy: 0.9561
/usr/local/lib/python3.10/dist-packages/sklearn/neighbors/_classification.py:239: DataConversionWarning: A column-vector y was passed wh
return self._fit(X, y)
```

### Performance Analysis:

- **Random Forest** performed the best with an accuracy of **97.37%**, thanks to its ensemble approach, which reduces overfitting and improves robustness.
- **Logistic Regression** and **SVM** achieved high accuracies of **95.61%** and **96.49%**, respectively, making them effective for binary classification with clear separations.
- **k-NN** had an accuracy of **92.98%**, slightly lower due to its reliance on computational power and performance with high-dimensional data.
- **Decision Tree** was the lowest with **92.11%**, likely due to overfitting, which can be mitigated with proper tuning.

```
# Dictionary to store the models
```

```
models = {
    "Logistic Regression": log_reg,
    "Decision Tree": dt_clf,
    "Random Forest": rf_clf,
    "SVM": svm_clf,
    "k-NN": knn_clf
}
```

```
# Dictionary to store the accuracy scores
```

```
accuracy_scores = {}
```

```
# Train each model and calculate accuracy
```

```
for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    accuracy_scores[name] = accuracy_score(y_test, y_pred)
```

```
# Print the accuracy of each model
```

```
for name, score in accuracy_scores.items():
    print(f"{name} Accuracy: {score:.4f}")
```

```
# Identify the best and worst performing models
```

```
best_model = max(accuracy_scores, key=accuracy_scores.get)
worst_model = min(accuracy_scores, key=accuracy_scores.get)
```

```
print(f"\nBest performing model: {best_model} with accuracy: {accuracy_scores[best_model]:.4f}")
print(f"Worst performing model: {worst_model} with accuracy: {accuracy_scores[worst_model]:.4f}")

/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1408: DataConversionWarning: A column-vector y was passed when a 1d
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:1389: DataConversionWarning: A column-vector y was passed when a 1d array was ex
return fit_method(estimator, *args, **kwargs)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1408: DataConversionWarning: A column-vector y was passed when a 1d
y = column_or_1d(y, warn=True)
Logistic Regression Accuracy: 0.9561
Decision Tree Accuracy: 0.9474
Random Forest Accuracy: 0.9649
SVM Accuracy: 0.9561
k-NN Accuracy: 0.9561

Best performing model: Random Forest with accuracy: 0.9649
Worst performing model: Decision Tree with accuracy: 0.9474
/usr/local/lib/python3.10/dist-packages/sklearn/neighbors/_classification.py:239: DataConversionWarning: A column-vector y was passed wf
return self._fit(X, y)
```

### Conclusion:

- **Best Model: Random Forest** achieved the highest accuracy, thanks to its ensemble approach that reduces overfitting and handles complexity well.
- **Worst Model: Decision Tree** performed the worst, likely due to overfitting, but can improve with proper tuning.

Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.